

# Vertx

Rafael Merino García

@imrafaelmerino

2021

# Table of Contents

- 1 Goals
- 2 History
- 3 Vertx
- 4 vertx-effect: where Vertx meet FP
- 5 hands-on

# Goals

# Goals (I)

- How to use **Vertx** following the **Erlang** philosophy (**message passing**)

# Goals (I)

- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and

# Goals (I)

- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and
  - How do you get programs to run faster **without changing a single line?**

# Goals (I)

- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and
  - How do you get programs to run faster **without changing a single line**?
    - Increasing CPU frequency?

- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and
  - How do you get programs to run faster **without changing a single line**?
    - Increasing CPU frequency?
    - Increasing number of cores?



- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and
  - How do you get programs to run faster **without changing a single line**?
    - Increasing CPU frequency?
    - Increasing number of cores?
    - Increasing number of machines?

- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and
  - How do you get programs to run faster **without changing a single line**?
    - Increasing CPU frequency?
    - Increasing number of cores?
    - Increasing number of machines?
- Understand better functional and reactive programming

# Goals (I)

- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and
  - How do you get programs to run faster **without changing a single line**?
    - Increasing CPU frequency?
    - Increasing number of cores?
    - Increasing number of machines?
- Understand better functional and reactive programming
  - FP shines dealing with **effects**

# Goals (I)

- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and
  - How do you get programs to run faster **without changing a single line**?
    - Increasing CPU frequency?
    - Increasing number of cores?
    - Increasing number of machines?
- Understand better functional and reactive programming
  - FP shines dealing with **effects**
  - Is your app really reactive?

- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and
  - How do you get programs to run faster **without changing a single line**?
    - Increasing CPU frequency?
    - Increasing number of cores?
    - Increasing number of machines?
- Understand better functional and reactive programming
  - FP shines dealing with **effects**
  - Is your app really reactive?
    - What's your first integration test?

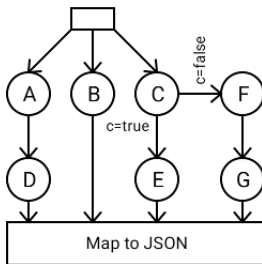
- How to use **Vertx** following the **Erlang** philosophy (**message passing**)
  - Why? **Low coupling** and
  - How do you get programs to run faster **without changing a single line**?
    - Increasing CPU frequency?
    - Increasing number of cores?
    - Increasing number of machines?
- Understand better functional and reactive programming
  - FP shines dealing with **effects**
  - Is your app really reactive?
    - What's your first integration test?
    - **Failures are just data**

## Goals (II)

- Develop complex services in Vertx **and not die trying** (async programming is still hard nowadays)

## Goals (II)


- Develop complex services in Vertx **and not die trying** (async programming is still hard nowadays)
- A, B and C are http requests or database calls executed **in parallel**





## Goals (II)

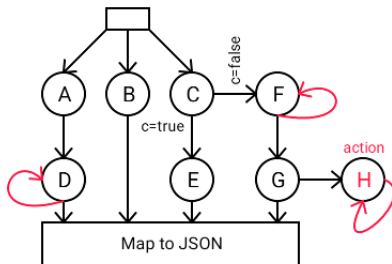
- Develop complex services in Vertx **and not die trying** (async programming is still hard nowadays)
- A, B and C are http requests or database calls executed **in parallel**



- I didn't mention, but I was thinking of **just one line of code**

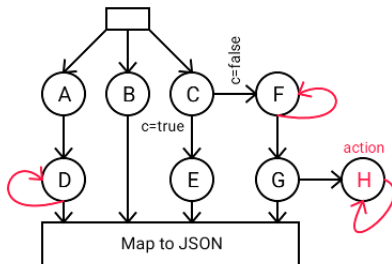
## Goals (III)

- and if you retry some ops under certain errors (connection timeouts, network partitions etc) or want to add some action without waiting for the response, like sending an email (H)



## Goals (III)

- and if you retry some ops under certain errors (connection timeouts, network partitions etc) or want to add some action without waiting for the response, like sending an email (H)



- I'm still thinking of just one line of code

# History

## Actor Model of Computation: Scalable Robust Information Systems

Carl Hewitt

*This article is dedicated to Alonzo Church and Dana Scott.*

The Actor Model is a mathematical theory that treats “*Actors*” as the universal primitives of digital computation.

Hypothesis:<sup>i</sup> **All physically possible computation can be directly implemented using Actors.**

The model has been used both as a framework for a theoretical understanding of concurrency, and as the theoretical basis for several practical implementations of concurrent systems. The advent of massive concurrency through client-cloud computing and many-core computer architectures has galvanized interest in the Actor Model.

**Message passing using types is the foundation of system communication:**

- Messages are the unit of communication<sup>1</sup>

- **Alan Kay**, one of its creator, **coined the term Object Oriented Programming** but...

- **Alan Kay**, one of its creator, **coined the term Object Oriented Programming** but...
- “I’m sorry that I long ago coined the term “objects” for this topic because it gets many people to focus on the lesser idea. **The big idea is messaging.**”

- **Alan Kay**, one of its creator, **coined the term Object Oriented Programming** but...
- "I'm sorry that I long ago coined the term "objects" for this topic because it gets many people to focus on the lesser idea. **The big idea is messaging.**"
- "I made up the term 'object-oriented', and I can tell you **I didn't have C++ in mind.**"



- **Alan Kay**, one of its creator, **coined the term Object Oriented Programming** but...
- "I'm sorry that I long ago coined the term "objects" for this topic because it gets many people to focus on the lesser idea. **The big idea is messaging.**"
- "I made up the term 'object-oriented', and I can tell you **I didn't have C++ in mind.**"
- Download Pharo and be blown away!

# Erlang (1986)

- Open Source in 1998

# Erlang (1986)

- Open Source in 1998
- Functional programming language created **to build fault-tolerant and scalable distributed systems in Ericsson**

# Erlang (1986)

- Open Source in 1998
- Functional programming language created **to build fault-tolerant and scalable distributed systems in Ericsson**
- Isolated processes that **interact sending messages** (10 servers handling 1 req instead of 1 server handling 10 req)

- Open Source in 1998
- Functional programming language created **to build fault-tolerant and scalable distributed systems in Ericsson**
- Isolated processes that **interact sending messages** (10 servers handling 1 req instead of 1 server handling 10 req)
- Keep calm and **let it crash** (avoid defensive programming and don't make things worse)

- Open Source in 1998
- Functional programming language created **to build fault-tolerant and scalable distributed systems in Ericsson**
- Isolated processes that **interact sending messages** (10 servers handling 1 req instead of 1 server handling 10 req)
- Keep calm and **let it crash** (avoid defensive programming and don't make things worse)
- WhatsApp and WeChat are implemented in Erlang!

# Influential people



# Vertx



- Understand what message passing means. **Adoption of a new paradigm is a complex and slow process that has to hurt.**

# Recomendations (my own)

- Understand what message passing means. **Adoption of a new paradigm is a complex and slow process that has to hurt.**
- EOOP just doesn't fit (**don't use Spring or try to avoid it**)

# Recomendations (my own)

- Understand what message passing means. **Adoption of a new paradigm is a complex and slow process that has to hurt.**
- EOOP just doesn't fit (**don't use Spring or try to avoid it**)
- Dependency injection == Coupling

# Recomendations (my own)

- Understand what message passing means. **Adoption of a new paradigm is a complex and slow process that has to hurt.**
- EOOP just doesn't fit (**don't use Spring or try to avoid it**)
- Dependency injection == Coupling
- **Have a plan to handle complexity.**

- **Everything** is a Verticle (minimum execution unit)

- **Everything** is a Verticle (minimum execution unit)
- computation | storage | communication

# Vertices

- **Everything** is a Verticle (minimum execution unit)
- computation | storage | communication
- Can create and kill other vertices

# Vertices

- **Everything** is a Verticle (minimum execution unit)
- computation | storage | communication
- Can create and kill other verticles
- Verticles are strongly **isolated** (share no resources)



- **Everything** is a Verticle (minimum execution unit)
- computation | storage | communication
- Can create and kill other vertices
- Verticles are strongly **isolated** (share no resources)
- Verticles creation and destruction is a lightweight operation

- **Everything** is a Verticle (minimum execution unit)
- computation | storage | communication
- Can create and kill other vertices
- Verticles are strongly **isolated** (share no resources)
- Verticles creation and destruction is a lightweight operation
- Verticles have **unique** addresses

- **Everything** is a Verticle (minimum execution unit)
- computation | storage | communication
- Can create and kill other verticles
- Verticles are strongly **isolated** (share no resources)
- Verticles creation and destruction is a lightweight operation
- Verticles have **unique** addresses
- If you know the address of a verticle you can send it a message

- **Everything** is a Verticle (minimum execution unit)
- computation | storage | communication
- Can create and kill other verticles
- Verticles are strongly **isolated** (share no resources)
- Verticles creation and destruction is a lightweight operation
- Verticles have **unique** addresses
- If you know the address of a verticle you can send it a message
- Message passing is **the only way** for verticles to interact

- **Everything** is a Verticle (minimum execution unit)
- computation | storage | communication
- Can create and kill other vertices
- Verticles are strongly **isolated** (share no resources)
- Verticles creation and destruction is a lightweight operation
- Verticles have **unique** addresses
- If you know the address of a verticle you can send it a message
- Message passing is **the only way** for vertices to interact
- **The more vertices the better**

# Message passing

- Message passing is assumed to be **atomic** which means that a message is either delivered in its entirety or not at all

# Message passing

- Message passing is assumed to be **atomic** which means that a message is either delivered in its entirety or not at all
- Message passing between a pair of vertices is assumed to be ordered: **the messages will be received in the same order they were sent**

# Message passing

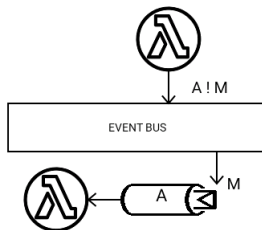
- Message passing is assumed to be **atomic** which means that a message is either delivered in its entirety or not at all
- Message passing between a pair of verticles is assumed to be ordered: **the messages will be received in the same order they were sent**
- Messages **should not contain references to data structures contained within verticles**—they should only contain constants and/or address



# Message passing

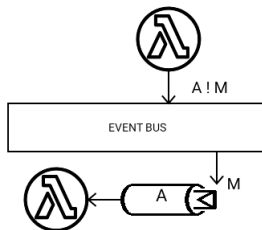
- Message passing is assumed to be **atomic** which means that a message is either delivered in its entirety or not at all
- Message passing between a pair of verticles is assumed to be ordered: **the messages will be received in the same order they were sent**
- Messages **should not contain references to data structures contained within verticles**—they should only contain constants and/or address
- **send and pray** semantics. We send the message and pray that it arrives

# Vertx model (I)



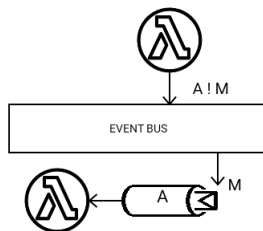
- The event bus is the nervous system of Vertx

# Vertx model (I)



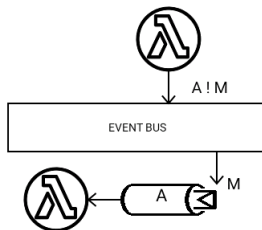
- The event bus is the nervous system of Vertx
- Distributed event bus: verticles **from different machines** can communicate with each other sending messages. Scale-out

# Vertx model (I)




- The event bus is the nervous system of Vertx
- Distributed event bus: verticles **from different machines** can communicate with each other sending messages. Scale-out
- It provides an API to:

# Vertx model (I)




- The event bus is the nervous system of Vertx
- Distributed event bus: verticles **from different machines** can communicate with each other sending messages. Scale-out
- It provides an API to:
  - deploy verticles listening on addresses

# Vertx model (I)




- The event bus is the nervous system of Vertx
- Distributed event bus: verticles **from different machines** can communicate with each other sending messages. Scale-out
- It provides an API to:
  - deploy verticles listening on addresses
  - **send and publish** messages to addresses: point-to-point and pub/sub

# Vertx model (I)



- The event bus is the nervous system of Vertx
- Distributed event bus: verticles **from different machines** can communicate with each other sending messages. Scale-out
- It provides an API to:
  - deploy verticles listening on addresses
  - **send and publish** messages to addresses: point-to-point and pub/sub
- **Verticles process ONE message at a time.** Synchronization is implemented this way


# Vertx model (II)



- Scale-up: deploying multiple instances of a verticle listening on the same address



# Vertx model (II)




- Scale-up: deploying multiple instances of a verticle listening on the same address
- Instance that receives the message is chosen using a non-strict round-robin algorithm

# Vertx model (III)

Imagine that a verticle **sends a message** to other verticle and has to **process the response**


- Programmatically, it's just a handler



# Vertx model (III)

Imagine that a verticle **sends a message** to other verticle and has to **process the response**

- Programmatically, it's just a handler
- but in practice it's just another verticle listening on a random address



# Verticles life cycle

- Verticles that are created during app bootstrap and never dye

# Verticles life cycle

- Verticles that are created during app bootstrap and never dye
- Verticles that are associated to the life cycle of an entity (user login and logout)

# Verticles life cycle

- Verticles that are created during app bootstrap and never dye
- Verticles that are associated to the life cycle of an entity (user login and logout)
- Verticles that are created to do computation and die after that. Why do we need an address then?

- String, null, bytes, Numbers, Boolean and Json (from Jackson)

# Messages

- String, null, bytes, Numbers, Boolean and Json (from Jackson)
- All supported messages have a MessageCodec associated



# Messages

- String, null, bytes, Numbers, Boolean and Json (from Jackson)
- All supported messages have a MessageCodec associated
- You can send your own messages if you create and register their codecs

# Messages

- String, null, bytes, Numbers, Boolean and Json (from Jackson)
- All supported messages have a MessageCodec associated
- You can send your own messages if you create and register their codecs
- The vertx Json is mutable. ¿Why is this a problem?

# Messages

- String, null, bytes, Numbers, Boolean and Json (from Jackson)
- All supported messages have a MessageCodec associated
- You can send your own messages if you create and register their codecs
- The vertex Json is mutable. ¿Why is this a problem?
  - Verticles have to be isolated

```
@Override
public JsonObject transform(JsonObject jsonObject) {
    return jsonObject.copy();
}
```

# Messages

- String, null, bytes, Numbers, Boolean and Json (from Jackson)
- All supported messages have a MessageCodec associated
- You can send your own messages if you create and register their codecs
- The vertex Json is mutable. ¿Why is this a problem?
  - Vertices have to be isolated

```
@Override
public JsonObject transform(JsonObject jsonObject) {
    return jsonObject.copy();
}
```

- [imrafaelmerino/json-values](#) is a better alternative: it's persistent and provides a better api

```
@Override
public JsObj transform(final JsObj obj) {
    return obj;
}
```

# Threading model

- Two type of threads: **event loops and workers**

# Threading model

- Two type of threads: **event loops and workers**
- When deploying a verticle you have to specify if it will be executed by a worker or an event loop

# Threading model

- Two type of threads: **event loops and workers**
- When deploying a verticle you have to specify if it will be executed by a worker or an event loop
- computationally intensive or blocking tasks  $\Rightarrow$  worker

# Threading model

- Two type of threads: **event loops and workers**
- When deploying a verticle you have to specify if it will be executed by a worker or an event loop
- computationally intensive or blocking tasks  $\Rightarrow$  worker
- never block an event loop



# Threading model

- Two type of threads: **event loops and workers**
- When deploying a verticle you have to specify if it will be executed by a worker or an event loop
- computationally intensive or blocking tasks  $\Rightarrow$  worker
- never block an event loop
- Size of pools by default

# Threading model

- Two type of threads: **event loops and workers**
- When deploying a verticle you have to specify if it will be executed by a worker or an event loop
- computationally intensive or blocking tasks  $\Rightarrow$  worker
- never block an event loop
- Size of pools by default
  - event loops: 2 x processors

# Threading model

- Two type of threads: **event loops and workers**
- When deploying a verticle you have to specify if it will be executed by a worker or an event loop
- computationally intensive or blocking tasks  $\Rightarrow$  worker
- never block an event loop
- Size of pools by default
  - event loops: 2 x processors
  - workers: 20

# Threading model

- Two type of threads: **event loops and workers**
- When deploying a verticle you have to specify if it will be executed by a worker or an event loop
- computationally intensive or blocking tasks  $\Rightarrow$  worker
- never block an event loop
- Size of pools by default
  - event loops: 2 x processors
  - workers: 20
  - What's the right size?

# Threading model

- Two type of threads: **event loops and workers**
- When deploying a verticle you have to specify if it will be executed by a worker or an event loop
- computationally intensive or blocking tasks  $\Rightarrow$  worker
- never block an event loop
- Size of pools by default
  - event loops: 2 x processors
  - workers: 20
  - What's the right size?
- What about green threads and project Loom?

vertx-effect: where Vertx meet FP

# Val and $\lambda$

```
import java.util.function.Supplier
import java.util.function.Function
import io.vertx.core.Future

public interface Val<O> extends Supplier<Future<O>> {...}

public interface  $\lambda$ <I,O> extends Function<I, Val<O>> {...}
```

- Val is lazy. It describes an asynchronous effect

# Val and $\lambda$

```
import java.util.function.Supplier
import java.util.function.Function
import io.vertx.core.Future

public interface Val<O> extends Supplier<Future<O>> {...}

public interface  $\lambda$ <I,O> extends Function<I, Val<O>> {...}
```

- Val is **lazy**. It **describes** an asynchronous effect
- The types **I** and **O** represent messages sent to the Event Bus



# Val and $\lambda$

```
import java.util.function.Supplier
import java.util.function.Function
import io.vertx.core.Future

public interface Val<O> extends Supplier<Future<O>> {...}

public interface  $\lambda$ <I,O> extends Function<I, Val<O>> {...}
```

- Val is **lazy**. It **describes** an asynchronous effect
- The types **I** and **O** represent messages sent to the Event Bus
- If they are not supported by Vertx

# Val and $\lambda$

```
import java.util.function.Supplier
import java.util.function.Function
import io.vertx.core.Future

public interface Val<O> extends Supplier<Future<O>> {...}

public interface  $\lambda$ <I,O> extends Function<I, Val<O>> {...}
```

- Val is **lazy**. It **describes** an asynchronous effect
- The types **I** and **O** represent messages sent to the Event Bus
- If they are not supported by Vertx
  - Implement and register a **MessageCodec** for them

hands-on



Michel Goossens, Frank Mittelbach, and Alexander Samarin. *The L<sup>A</sup>T<sub>E</sub>X Companion*. Addison-Wesley, Reading, Massachusetts, 1993.



Albert Einstein. *Zur Elektrodynamik bewegter Körper*. (German) [*On the electrodynamics of moving bodies*]. *Annalen der Physik*, 322(10):891–921, 1905.



Knuth: Computers and Typesetting,  
<http://www-cs-faculty.stanford.edu/~uno/abcde.html>